

# WATER WELLS IN TANZANIA.

## Overview and Problem Understanding.

After careful evaluation , it has been noted that most wells in tanzania have been experiencing trouble and are faulty. A team of 5 was given data and told to help Navesco Company (a water company in Tanzania) to come up with eloquent data and machine understanding to come up with different ways in which new and improved wells would help Tanzanians stop experiencing dcrougt yet water is sufficient in the country

We were able to get data from a reliable source and it will inturn be of help to us coming up with different models with different results to help in the aid of this project.

Objectives of this project include the following :

1. Build a classifier to predict the condition of a water well and pitch to an NGO focused on locating wells needing repair.

## DATA UNDERSTANDING

The data provided was divided into three different stages and levels which include the following :

1. Testing data values
2. Training data values
3. Testing labels to aid in model creation.

We are required to join the data or use the data just as it is to come up with two or more different models and explanations for the problem .

## DATA EXPLORATION.

We will begin by firstly, importing necessary libraries that will allow us to open the data and get to know what is present in our files .


```
In [1]: # import necessary libraries.  
import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
import warnings  
warnings.filterwarnings("ignore")  
  
#import train and test data libraries  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import LabelEncoder  
from sklearn.preprocessing import StandardScaler  
from sklearn.impute import SimpleImputer  
from sklearn.model_selection import GridSearchCV, cross_validate  
from sklearn.model_selection import cross_val_predict, cross_val_score,  
from sklearn.pipeline import Pipeline  
from imblearn.over_sampling import SMOTE  
  
#metrics for baseline dataset  
from sklearn.metrics import confusion_matrix, accuracy_score, f1_score,  
from sklearn.metrics import ConfusionMatrixDisplay, classification_report  
from sklearn.metrics import roc_curve, auc, roc_auc_score  
  
# models to account for  
from sklearn.linear_model import LogisticRegression  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.ensemble import BaggingClassifier, RandomForestClassifier  
from sklearn.neighbors import KNeighborsClassifier  
  
# reduce data memory  
from scipy.sparse import csr_matrix
```


```
In [2]: # open our test set data i.e all files and see the data info within the
data1 = pd.read_csv("training_set_values.csv")
pd.DataFrame(data1)
data1
```

Out[2]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	.
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	.
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	.
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	.
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	.
...	...	...	...	...	...	...	...	...
59395	60739	10.0	2013-05-03	Germany Republi	1210	CES	37.169807	.
59396	27263	4700.0	2011-05-07	Cefa-njombe	1212	Cefa	35.249991	.
59397	37057	0.0	2011-04-11	NaN	0	NaN	34.017087	.
59398	31282	0.0	2011-03-08	Malec	0	Musa	35.861315	.
59399	26348	0.0	2011-03-23	World Bank	191	World	38.104048	.

59400 rows × 40 columns



```
In [3]:  #open the test variables (i.e labels)  
data2 = pd.read_csv('training_set_labels.csv')  
data2
```

```
Out[3]:
```

	id	status_group
0	69572	functional
1	8776	functional
2	34310	functional
3	67743	non functional
4	19728	functional
...	...	...
59395	60739	functional
59396	27263	functional
59397	37057	functional
59398	31282	functional
59399	26348	functional

59400 rows × 2 columns


Using the data above , we can now check the counts present for all values i.e functional,non-functional and spoilt wells .

```
In [4]: #join the columns with the training data that resonates.  
data2.value_counts()  
  
df = pd.merge(data1, data2, how = 'left', on='id')  
df
```

Out[4]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	.
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	.
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	.
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	.
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	.
...	...	...	...	...	...	...	...	...
59395	60739	10.0	2013-05-03	Germany Republi	1210	CES	37.169807	.
59396	27263	4700.0	2011-05-07	Cefa-njombe	1212	Cefa	35.249991	.
59397	37057	0.0	2011-04-11	NaN	0	NaN	34.017087	.
59398	31282	0.0	2011-03-08	Malec	0	Musa	35.861315	.
59399	26348	0.0	2011-03-23	World Bank	191	World	38.104048	.

59400 rows × 41 columns



In [5]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 41 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     59400 non-null  int64
1   amount_tsh                           59400 non-null  float64
2   date_recorded                         59400 non-null  object
3   funder                                55763 non-null  object
4   gps_height                            59400 non-null  int64
5   installer                             55745 non-null  object
6   longitude                             59400 non-null  float64
7   latitude                             59400 non-null  float64
8   wpt_name                              59398 non-null  object
9   num_private                           59400 non-null  int64
10  basin                                 59400 non-null  object
11  subvillage                            59029 non-null  object
12  region                                59400 non-null  object
13  region_code                           59400 non-null  int64
14  district_code                         59400 non-null  int64
15  lga                                    59400 non-null  object
16  ward                                  59400 non-null  object
17  population                             59400 non-null  int64
18  public_meeting                        56066 non-null  object
19  recorded_by                           59400 non-null  object
20  scheme_management                     55522 non-null  object
21  scheme_name                           30590 non-null  object
22  permit                                56344 non-null  object
23  construction_year                     59400 non-null  int64
24  extraction_type                       59400 non-null  object
25  extraction_type_group                  59400 non-null  object
26  extraction_type_class                  59400 non-null  object
27  management                             59400 non-null  object
28  management_group                       59400 non-null  object
29  payment                                59400 non-null  object
30  payment_type                           59400 non-null  object
31  water_quality                          59400 non-null  object
32  quality_group                          59400 non-null  object
33  quantity                               59400 non-null  object
34  quantity_group                         59400 non-null  object
35  source                                 59400 non-null  object
36  source_type                           59400 non-null  object
37  source_class                           59400 non-null  object
38  waterpoint_type                        59400 non-null  object
39  waterpoint_type_group                  59400 non-null  object
40  status_group                           59400 non-null  object
dtypes: float64(3), int64(7), object(31)
memory usage: 18.6+ MB

```

In [6]: `df1 = df.copy()`

```
In [7]: df1 = df1.sample(frac=0.4, random_state=42)
df1['permit'] = df1['permit'].dropna()
df1
```

```
Out[7]:
```

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitud
<b>2980</b>	37098	0.0	2012-10-09	Rural Water Supply And Sanitat	0	DWE	31.98561
<b>5246</b>	14530	0.0	2012-11-03	Halmashauri Ya Manispa Tabora	0	Halmashauri ya manispa tabora	32.83281
<b>22659</b>	62607	10.0	2013-02-25	Siter Fransis	1675	DWE	35.48821
<b>39888</b>	46053	0.0	2011-08-13	Kkkt	0	KKKT	33.14081
<b>13361</b>	47083	50.0	2013-02-08	Wateraid	1109	SEMA	34.21701
...	...	...	...	...	...	...	...
<b>46912</b>	50526	0.0	2013-02-19	Adra	1564	Adra	35.52211
<b>41173</b>	67498	0.0	2013-02-17	Netherlands	0	DWE	33.68451
<b>37745</b>	71958	10.0	2011-02-27	Dwe	-2	DWE	38.64751
<b>14756</b>	52297	0.0	2013-03-26	Losaa-kia Water Supply	1231	Losaa-Kia water supp	37.11441
<b>55518</b>	34993	0.0	2011-08-02	Hifab	0	Hesawa	33.34191

23760 rows × 41 columns


## DATA PROCESSING

Using the dataset created , we can go ahead and begin removal of duplicates and null values and even columns that may not be of help especially with our objectives in play .

```
In [8]: #checking for duplicates
df1.duplicated().sum()
```

```
Out[8]: 0
```

our observation is that no data is duplicated and all is well with our dataset.

```
In [9]:  #check for null values  
df1.isna().sum()
```

```
Out[9]: id                                0  
amount_tsh                               0  
date_recorded                            0  
funder                                   1485  
gps_height                               0  
installer                               1492  
longitude                                0  
latitude                                 0  
wpt_name                                 1  
num_private                              0  
basin                                    0  
subvillage                              143  
region                                   0  
region_code                             0  
district_code                           0  
lga                                       0  
ward                                      0  
population                              0  
public_meeting                          1319  
recorded_by                              0  
scheme_management                       1553  
scheme_name                             11516  
permit                                  1226  
construction_year                       0  
extraction_type                         0  
extraction_type_group                   0  
extraction_type_class                   0  
management                             0  
management_group                       0  
payment                                 0  
payment_type                            0  
water_quality                           0  
quality_group                           0  
quantity                                0  
quantity_group                          0  
source                                   0  
source_type                             0  
source_class                            0  
waterpoint_type                         0  
waterpoint_type_group                   0  
status_group                            0  
dtype: int64
```

Using the details above, we find out that most columns are not going to be of help especially with the objectives present .So inorder to have clear columnns , we would be required to drop the unwanted columnns and also remove any null values from the needed columns.



In [10]: `#check the number of columns`  
`df1.columns`

Out[10]: Index(['id', 'amount\_tsh', 'date\_recorded', 'funder', 'gps\_height',  
 'installer', 'longitude', 'latitude', 'wpt\_name', 'num\_private',  
 'basin', 'subvillage', 'region', 'region\_code', 'district\_code', 'lga',  
 'ward', 'population', 'public\_meeting', 'recorded\_by',  
 'scheme\_management', 'scheme\_name', 'permit', 'construction\_year',  
 'extraction\_type', 'extraction\_type\_group', 'extraction\_type\_class',  
 'management', 'management\_group', 'payment', 'payment\_type',  
 'water\_quality', 'quality\_group', 'quantity', 'quantity\_group',  
 'source', 'source\_type', 'source\_class', 'waterpoint\_type',  
 'waterpoint\_type\_group', 'status\_group'],  
 dtype='object')

Due to the presence of many unrequired columns , we would need to impute them and drop the unwanted columns

In [11]: `# differentiate the numeric columns to the categorical columns`  
`df1.select_dtypes(include=['object']).columns`

Out[11]: Index(['date\_recorded', 'funder', 'installer', 'wpt\_name', 'basin',  
 'subvillage', 'region', 'lga', 'ward', 'public\_meeting', 'recorded\_by',  
 'scheme\_management', 'scheme\_name', 'permit', 'extraction\_type',  
 'extraction\_type\_group', 'extraction\_type\_class', 'management',  
 'management\_group', 'payment', 'payment\_type', 'water\_quality',  
 'quality\_group', 'quantity', 'quantity\_group', 'source', 'source\_type',  
 'source\_class', 'waterpoint\_type', 'waterpoint\_type\_group',  
 'status\_group'],  
 dtype='object')

In [12]: `df1.select_dtypes(include = ['int64', 'float64']).columns`

Out[12]: Index(['id', 'amount\_tsh', 'gps\_height', 'longitude', 'latitude',  
 'num\_private', 'region\_code', 'district\_code', 'population',  
 'construction\_year'],  
 dtype='object')

In [13]: `# drop the unrequired columns to avoid data duplication`  
`columns_to_drop = ['date_recorded', 'funder', 'wpt_name', 'subvillage',  
 'ward', 'recorded_by', 'scheme_name', 'extraction_type',  
 'extraction_type_group', 'management', 'payment', 'quality_group',  
 'quantity', 'source', 'source_type', 'waterpoint_type', 'num_private',  
 'region_code', 'district_code']`  
`df1 = df1.drop(columns_to_drop, axis=1)`

```
In [14]: df2 = df1.copy().dropna()
df2.head()
```

```
Out[14]:
```

	id	amount_tsh	gps_height	installer	longitude	latitude	basin	
2980	37098	0.0	0	DWE	31.985658	-3.596360	Lake Tanganyika	Str
5246	14530	0.0	0	Halmashauri ya manispa tabora	32.832815	-4.944937	Lake Tanganyika	
22659	62607	10.0	1675	DWE	35.488289	-4.242048	Internal	I
39888	46053	0.0	0	KKKT	33.140828	-9.059386	Lake Rukwa	
13361	47083	50.0	1109	SEMA	34.217077	-4.430529	Internal	

5 rows × 21 columns

## EDA ANALYSIS

### UNIVARIATE ANALYSIS

Now that the data is all cleaned and without any null values, we will begin with the analysis section of the project. Firstly, we begin with different singular analysis between different columns and relations.

```
In [15]: df2.columns
```

```
Out[15]: Index(['id', 'amount_tsh', 'gps_height', 'installer', 'longitude', 'latitude',
               'basin', 'region', 'population', 'public_meeting', 'scheme_management',
               'permit', 'construction_year', 'extraction_type_class',
               'management_group', 'payment_type', 'water_quality', 'quantity_group',
               'source_class', 'waterpoint_type_group', 'status_group'],
              dtype='object')
```

```
In [16]: #Numerical analysis('amount_tsh')
print(df2['amount_tsh'].describe())
```

```
count    19490.000000
mean       362.575947
std       2897.612079
min         0.000000
25%         0.000000
50%         0.000000
75%        50.000000
max      200000.000000
Name: amount_tsh, dtype: float64
```

In [17]: `#categorical data ('payment_type')`  
`print(df2['payment_type'].describe())`


```
count          19490
unique           7
top      never pay
freq           8365
Name: payment_type, dtype: object
```

In [18]: `#limiting the number of categories present in our training dataset or ou`  
`# Limit to the top 4 most frequent categories`  
`top_categories = df2['waterpoint_type_group'].value_counts().nlargest(4)`  
  
`# Filter the data to only include the top categories`  
`df_filtered = df2[df2['waterpoint_type_group'].isin(top_categories)]`  
`df_filtered`

Out[18]:

	id	amount_tsh	gps_height	installer	longitude	latitude	basin	
<b>2980</b>	37098	0.0	0	DWE	31.985658	-3.596360	Lake Tanganyika	SI
<b>5246</b>	14530	0.0	0	Halmashauri ya manispa tabora	32.832815	-4.944937	Lake Tanganyika	
<b>22659</b>	62607	10.0	1675	DWE	35.488289	-4.242048	Internal	
<b>39888</b>	46053	0.0	0	KKKT	33.140828	-9.059386	Lake Rukwa	
<b>13361</b>	47083	50.0	1109	SEMA	34.217077	-4.430529	Internal	
...	...	...	...	...	...	...	...	
<b>46912</b>	50526	0.0	1564	Adra	35.522134	-4.480657	Internal	
<b>41173</b>	67498	0.0	0	DWE	33.684507	-3.197790	Lake Victoria	SI
<b>37745</b>	71958	10.0	-2	DWE	38.647507	-7.869125	Rufiji	
<b>14756</b>	52297	0.0	1231	Losaa-Kia water supp	37.114494	-3.215041	Pangani	Kil
<b>55518</b>	34993	0.0	0	Hesawa	33.341961	-3.026684	Lake Victoria	

19455 rows × 21 columns

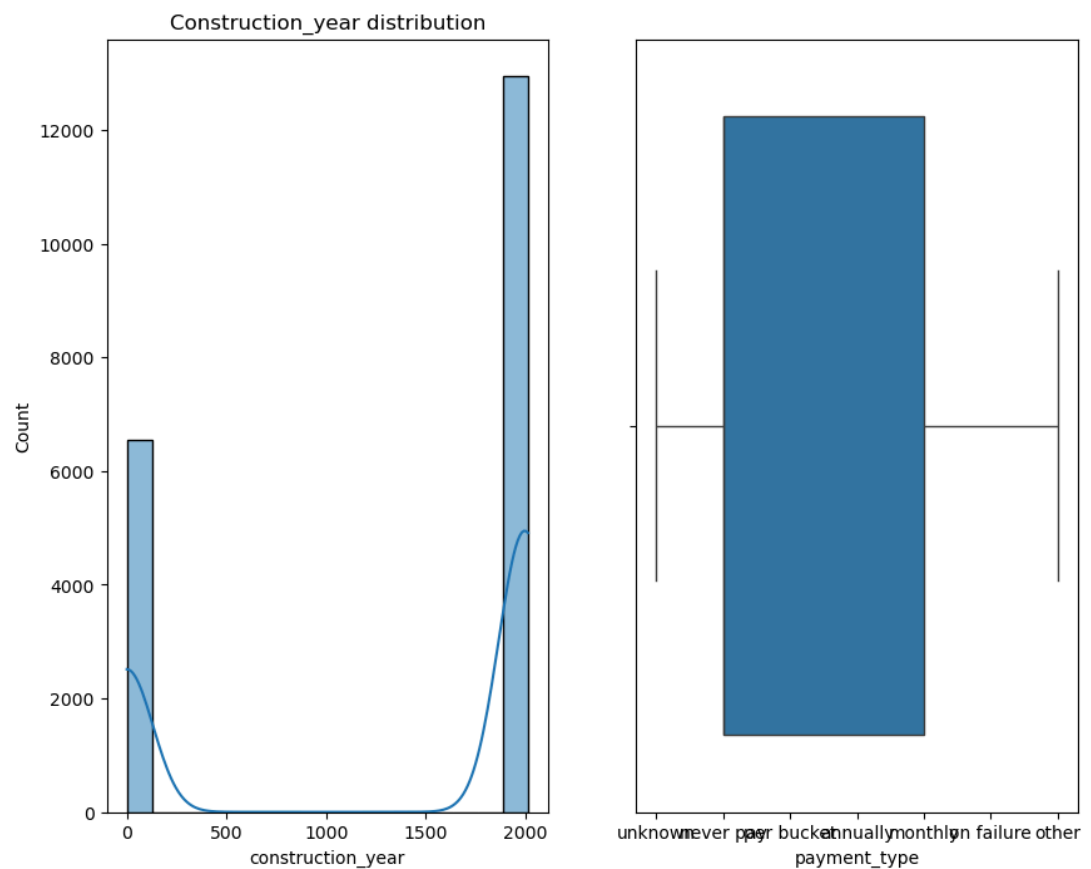


```
In [19]: #histogram for ('construction_year')
plt.figure(figsize=(10,8))

plt.subplot(1, 2, 1)
sns.histplot(df2['construction_year'], kde=True)
plt.title('Construction_year distribution')

#Boxplot for 'waterpoint type group'
plt.subplot(1, 2, 2)
sns.boxplot(x=df_filtered['payment_type'])
plt.title('')
```

Out[19]: Text(0.5, 1.0, '')



## BIVARIATE ANALYSIS

Now on to bivariate analysis where we analyse different relationships between two data types i.e 1. numerical vs categorical 2. numerical vs numerical 3. categorical vs categorical

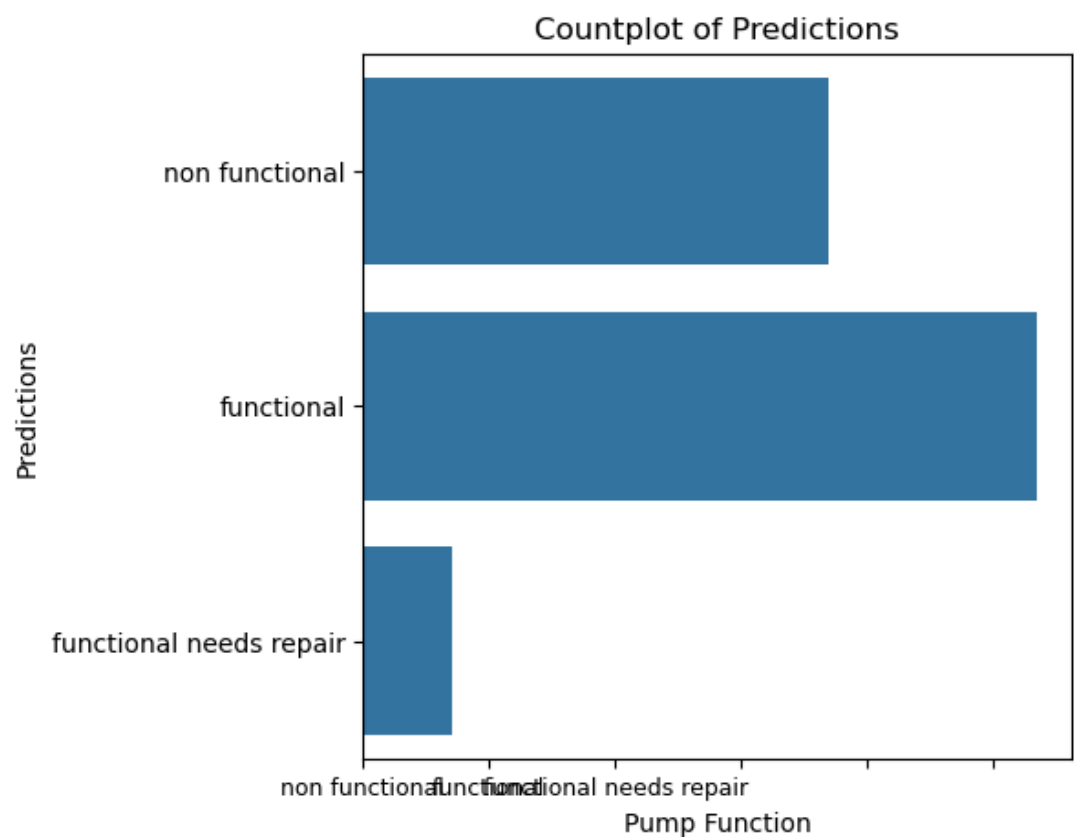
In [20]: `df2.describe()`

Out[20]:

	id	amount_tsh	gps_height	longitude	latitude	popul
<b>count</b>	19490.000000	19490.000000	19490.000000	19490.000000	1.949000e+04	19490.00
<b>mean</b>	37231.279785	362.575947	699.905798	34.163893	-5.679083e+00	176.71
<b>std</b>	21448.342293	2897.612079	701.891554	6.430768	2.872993e+00	503.05
<b>min</b>	3.000000	0.000000	-90.000000	0.000000	-1.158630e+01	0.00
<b>25%</b>	18561.250000	0.000000	0.000000	33.032037	-8.218093e+00	0.00
<b>50%</b>	37414.500000	0.000000	461.500000	35.197609	-4.985169e+00	30.00
<b>75%</b>	55686.750000	50.000000	1336.000000	37.326356	-3.330653e+00	200.00
<b>max</b>	74246.000000	200000.000000	2623.000000	40.322805	-2.000000e-08	30500.00

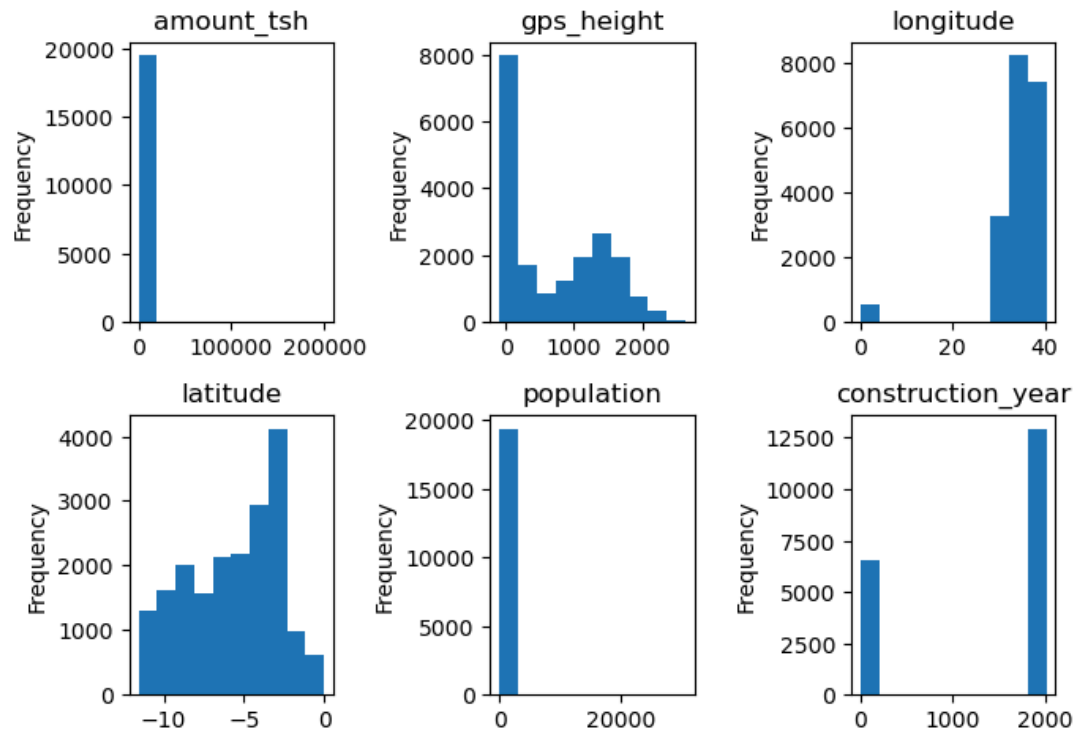
```
In [21]: # Plotting distribution of target variable.
fig, ax = plt.subplots(figsize=(5, 5))
sns.countplot(df2['status_group'])
x_labels = df2['status_group'].unique()

# Add Labels
plt.title('Countplot of Predictions')
plt.xlabel('Pump Function')
ax.set_xticklabels(x_labels, fontsize=9)
plt.ylabel('Predictions')
plt.show()
```

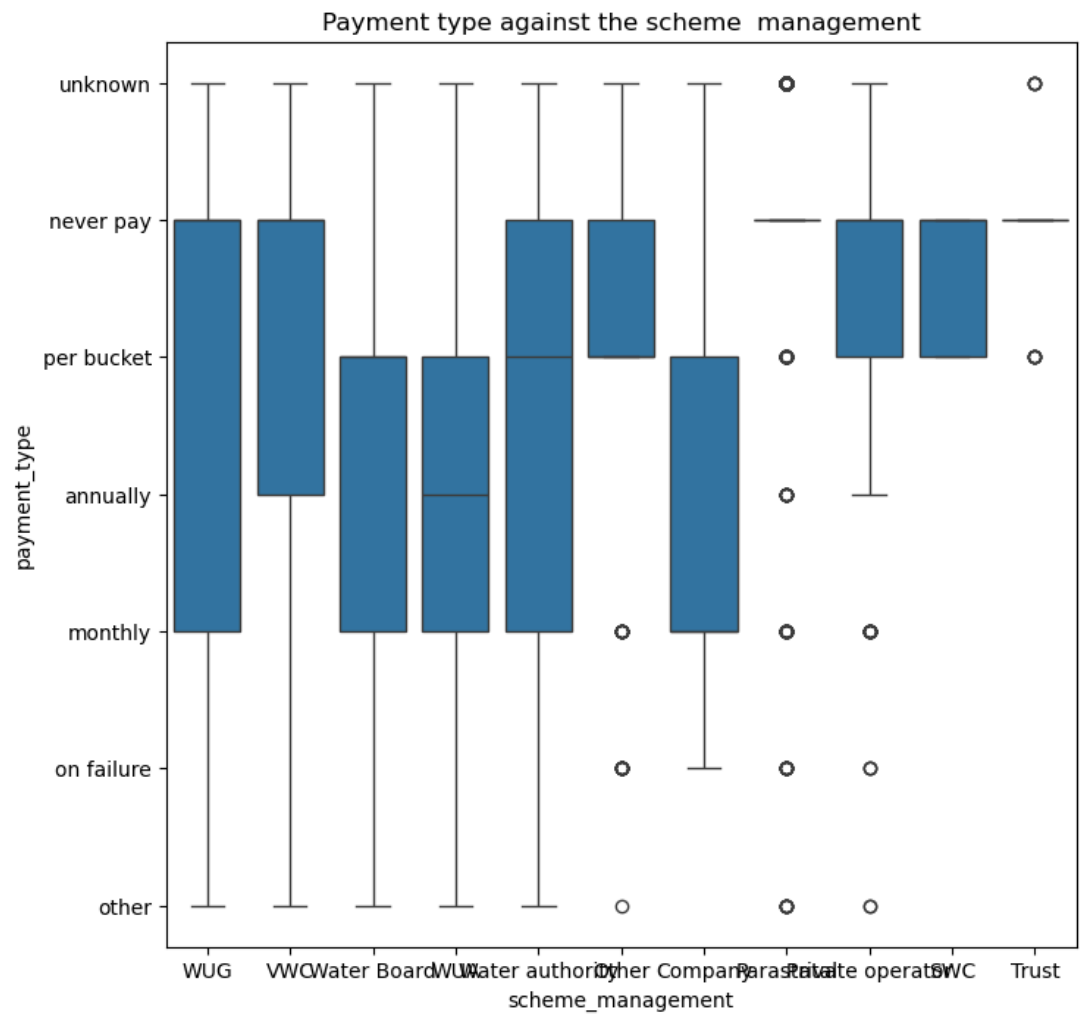



With this , we can now note that the functional wells are the ones with a high rate comared

```
In [22]: ▶ #for categorical data , we can use a histogram to represent
continuous = ['amount_tsh', 'gps_height', 'longitude', 'latitude', 'population']
fig = plt.figure(figsize=(7, 7))
for i, col in enumerate(continuous):
    ax = plt.subplot(3, 3, i+1)
    df2[col].plot(kind='hist', ax=ax, title=col)
plt.tight_layout()
plt.show()
```



```
In [23]: # Box plot for 'scheme_management' grouped by 'payment_type'
plt.figure(figsize=(8, 8))
sns.boxplot(x='scheme_management', y='payment_type', data=df_filtered)
plt.title('Payment type against the scheme management')
plt.show()
```



```
In [24]:  #RELATIONSHIP BETWEEN COMPANY FUNDING AND AMOUNT PROVIDED
grouped = df.groupby('installer')['amount_tsh'].describe()
print(grouped)
```

	count	mean	std	min	25%	50%	
75% \ installer							
-	3.0	0.000000	0.000000	0.0	0.0	0.0	
0.0							
0	777.0	243.885457	2658.671884	0.0	0.0	50.0	5
0.0							
A.D.B	1.0	0.000000	NaN	0.0	0.0	0.0	
0.0							
AAR	4.0	0.000000	0.000000	0.0	0.0	0.0	
0.0							
ABASIA	29.0	355.172414	222.944911	0.0	100.0	500.0	50
0.0							
...	...	...	...	...	...	...	
...							
water board	2.0	0.000000	0.000000	0.0	0.0	0.0	
0.0							
wizara ya maji	2.0	0.000000	0.000000	0.0	0.0	0.0	
0.0							
world	4.0	0.000000	0.000000	0.0	0.0	0.0	
0.0							
world banks	1.0	50.000000	NaN	50.0	50.0	50.0	5
0.0							
world vision	1.0	0.000000	NaN	0.0	0.0	0.0	
0.0							
	max						
installer							
-	0.0						
0	45000.0						
A.D.B	0.0						
AAR	0.0						
ABASIA	500.0						
...	...						
water board	0.0						
wizara ya maji	0.0						
world	0.0						
world banks	50.0						
world vision	0.0						
[2145 rows x 8 columns]							


In [ ]: 

## MULTIVARIATE ANALYSIS

After careful bivariate and univariate analysis , we can begin on multivariate to ensure complete analysis on the data present.



```
In [25]: #for the numerical details , we will use a correlation heatmap  
correlation_matrix = df_filtered[['construction_year', 'amount_tsh', 'pop  
plt.figure(figsize=(8, 8))  
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')  
plt.title('Correlation Heatmap for most numerical values')  
plt.show()
```

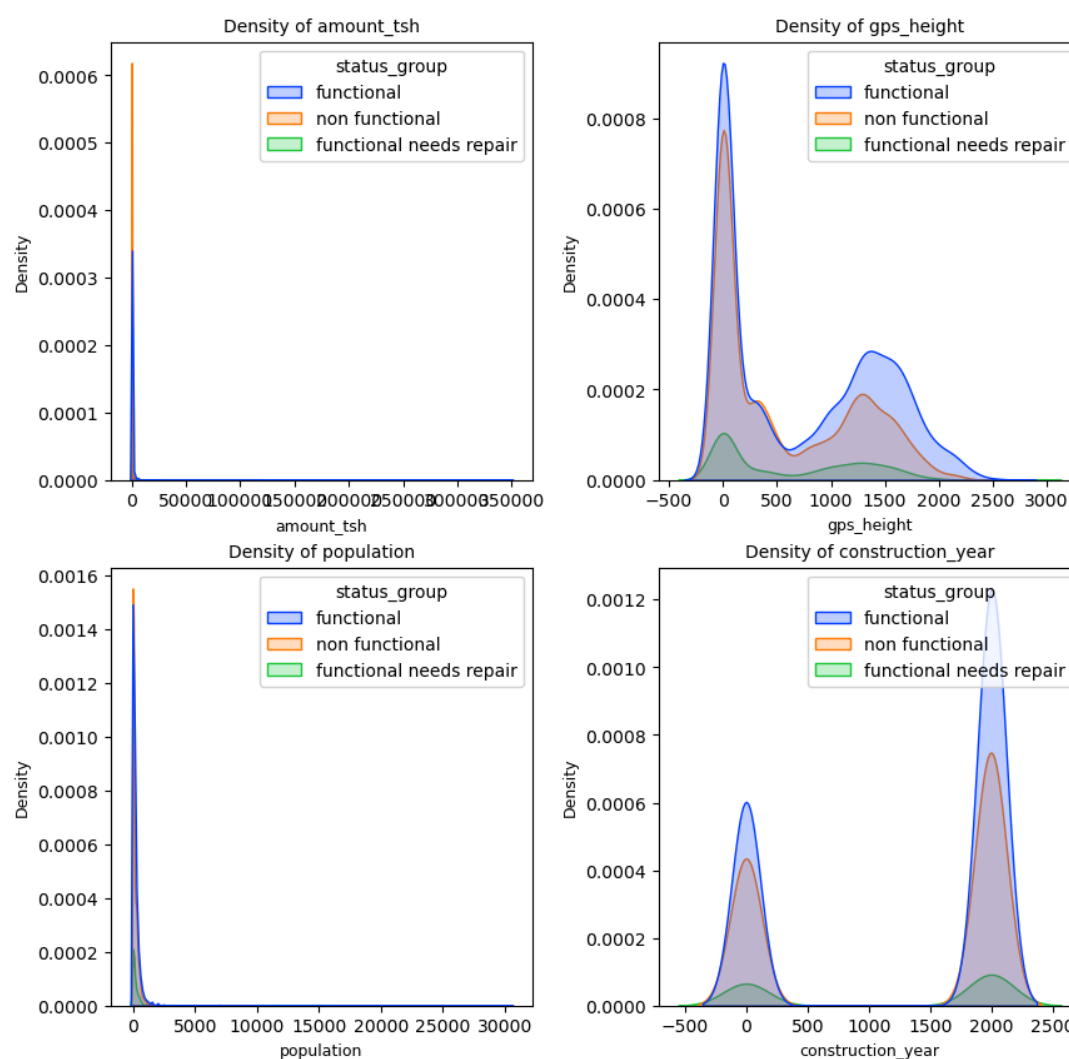


```
In [26]: # check the target variable against continuous data values
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(10, 10))
fig.subplots_adjust(hspace=0.2, wspace=0.3) # Increase hspace and wspace

# Create count bar plots for continuous variables
continuous = ['amount_tsh', 'gps_height',
              'population', 'construction_year']

for i, column in enumerate(continuous):
    row, col = i // 2, i % 2
    sns.kdeplot(data=df, x=column, hue='status_group', palette='bright',
                fill=True, ax=axes[row, col])
    axes[row, col].set_title(f'Density of {column}', fontsize=10)
    axes[row, col].set_xlabel(column, fontsize=9)
    axes[row, col].set_ylabel('Density', fontsize=9)

plt.show()
```



Now that we have been successful in data cleaning and eda analysis, we can now begin with the machine learning section .This involves now getting into the nitty gritty section of the project .Here we can evaluate the objectives and come up with different observations that will help us gain more insight into the data .

## FEATURE AND PREDICTOR

```
In [27]:  # choosing the target(status_group) and all others as the target variable
y = df2['status_group']

X = df2.drop('status_group', axis=1)

X.head()
```

```
Out[27]:
```

	id	amount_tsh	gps_height	installer	longitude	latitude	basin	
2980	37098	0.0	0	DWE	31.985658	-3.596360	Lake Tanganyika	Str
5246	14530	0.0	0	Halmashauri ya manispa tabora	32.832815	-4.944937	Lake Tanganyika	
22659	62607	10.0	1675	DWE	35.488289	-4.242048	Internal	I
39888	46053	0.0	0	KKKT	33.140828	-9.059386	Lake Rukwa	
13361	47083	50.0	1109	SEMA	34.217077	-4.430529	Internal	

## MODEL AND MACHINE LEARNING.

```
In [28]:  #convert categorical features and binary features
categ_col = ['installer', 'basin', 'region', 'scheme_management',
             'extraction_type_class', 'management_group', 'payment_type',
             'water_quality', 'quantity_group', 'source_class',
             'waterpoint_type_group']

cont_col = ['amount_tsh', 'gps_height', 'longitude', 'latitude', 'population']

binary_col = ['public_meeting', 'permit']
```

Using the grouped columns into either category, binary or count , we can now get dummies for the categorical columns present.

```
In [29]: #using label encoding to get the integral values for the categorical val
X = pd.get_dummies(X , columns = categ_col)
X['permit'] = X['permit'].dropna()
# Print X
X
```

Out[29]:

	id	amount_tsh	gps_height	longitude	latitude	population	public_meeting
<b>2980</b>	37098	0.0	0	31.985658	-3.596360	0	True
<b>5246</b>	14530	0.0	0	32.832815	-4.944937	0	True
<b>22659</b>	62607	10.0	1675	35.488289	-4.242048	148	True
<b>39888</b>	46053	0.0	0	33.140828	-9.059386	0	False
<b>13361</b>	47083	50.0	1109	34.217077	-4.430529	235	True
...	...	...	...	...	...	...	...
<b>46912</b>	50526	0.0	1564	35.522134	-4.480657	786	True
<b>41173</b>	67498	0.0	0	33.684507	-3.197790	0	True
<b>37745</b>	71958	10.0	-2	38.647507	-7.869125	1	True
<b>14756</b>	52297	0.0	1231	37.114494	-3.215041	1	True
<b>55518</b>	34993	0.0	0	33.341961	-3.026684	0	True

19490 rows × 1254 columns




In [30]: **X**

Out[30]:

	id	amount_tsh	gps_height	longitude	latitude	population	public_meeting
<b>2980</b>	37098	0.0	0	31.985658	-3.596360	0	True
<b>5246</b>	14530	0.0	0	32.832815	-4.944937	0	True
<b>22659</b>	62607	10.0	1675	35.488289	-4.242048	148	True
<b>39888</b>	46053	0.0	0	33.140828	-9.059386	0	False
<b>13361</b>	47083	50.0	1109	34.217077	-4.430529	235	True
...	...	...	...	...	...	...	...
<b>46912</b>	50526	0.0	1564	35.522134	-4.480657	786	True
<b>41173</b>	67498	0.0	0	33.684507	-3.197790	0	True
<b>37745</b>	71958	10.0	-2	38.647507	-7.869125	1	True
<b>14756</b>	52297	0.0	1231	37.114494	-3.215041	1	True
<b>55518</b>	34993	0.0	0	33.341961	-3.026684	0	True

19490 rows × 1254 columns



Now I will continue with train and test splitting

In [31]: `X_train , X_test , y_train ,y_test = train_test_split(X,y, test_size=0.3`

In [32]: `# Check for NaN values in the DataFrame  
print(X.isna().sum()) # Count of NaNs per column`

```
id          0
amount_tsh  0
gps_height  0
longitude   0
latitude    0
..
waterpoint_type_group_communal standpipe  0
waterpoint_type_group_dam                0
waterpoint_type_group_hand pump          0
waterpoint_type_group_improved spring    0
waterpoint_type_group_other              0
Length: 1254, dtype: int64
```

Using this we can now begin on model training

In [33]: `#check the training data information  
X_train.describe()`

Out[33]:

	id	amount_tsh	gps_height	longitude	latitude	popul
<b>count</b>	13643.000000	13643.000000	13643.000000	13643.000000	1.364300e+04	13643.00
<b>mean</b>	37204.715092	372.713626	696.928242	34.159379	-5.676880e+00	177.35
<b>std</b>	21375.562558	3159.961370	699.330855	6.483429	2.878076e+00	530.80
<b>min</b>	4.000000	0.000000	-90.000000	0.000000	-1.158630e+01	0.00
<b>25%</b>	18713.000000	0.000000	0.000000	33.037126	-8.223847e+00	0.00
<b>50%</b>	37251.000000	0.000000	457.000000	35.240377	-4.996336e+00	25.00
<b>75%</b>	55598.500000	50.000000	1332.000000	37.342572	-3.328403e+00	200.00
<b>max</b>	74238.000000	200000.000000	2585.000000	40.322625	-2.000000e-08	30500.00

I will have to scale it due to the incoherent values and scales present in our data and to also check the performance of the model

```
In [34]: ▶ #check for the metrics list as a global variable
metrics_list =[]

class Metrics:
    def model_score(self, model, y_true, y_pred):
        # Print classification report, accuracy, precision, recall, f1_score
        print(classification_report(y_true, y_pred))
        print("Overall accuracy score", accuracy_score(y_true, y_pred))
        print("Overall precision score", precision_score(y_true, y_pred, average='weighted'))
        print("Overall recall score", recall_score(y_true, y_pred, average='weighted'))
        print("Overall F1-score", f1_score(y_true, y_pred, average='weighted'))

        # Print a confusion matrix
        cnf_matrix = confusion_matrix(y_true, y_pred)
        disp = ConfusionMatrixDisplay(confusion_matrix=cnf_matrix, display=cnf_matrix)
        print('\nConfusion Matrix')
        return disp.plot()

    # Create a List of model metrics
    def get_metrics(self, model_name, model, y_true, y_pred): #y_test,
        metrics = {}
        metrics['model_name'] = model_name
        metrics['accuracy'] = accuracy_score(y_true, y_pred)
        metrics['f1 score'] = f1_score(y_true, y_pred, average='weighted')
        metrics['precision'] = precision_score(y_true, y_pred, average='weighted')
        metrics['recall'] = recall_score(y_true, y_pred, average='weighted')
        metrics_list.append(metrics)
        return metrics_list
```

## PIPELINE MODELLING

```
In [35]: ▶ # we can now move onto the next step of pipeline modelling
pipe_logistic = Pipeline([('ss', StandardScaler()),
                           ('lr', LogisticRegression(random_state=42))])

#pipeline classifier for Decision tree
pipe_tree = Pipeline([('ss', StandardScaler()),
                       ('tree', DecisionTreeClassifier(random_state=42))])

#pipeline classifier for the KNN
pipe_knn = Pipeline([('ss', StandardScaler()),
                      ('knn', KNeighborsClassifier)])
```

Now that we have already chosen the models to create using the original data above , we can now move to building the models one at a time .

## Regression Model

Here we can go ahead and instantiate the regression model and fit in our data and the model classifiers

```
In [36]: # Fitting the model
logistic_model = pipe_logistic.fit(X_train , y_train)

# print the accuracy levels for the test data
pipe_logistic.score(X_test,y_test)
```

Out[36]: 0.7400376261330597

```
In [37]: # Print model metrics
# Predict target
y_pred = logistic_model .predict(X_test)

# Create metrics object
score_metrics = Metrics()

# Print classification report, scores, and confusion matrix
score_metrics.model_score(logistic_model, y_test, y_pred)
```

	precision	recall	f1-score	support
functional	0.74	0.87	0.80	3197
functional needs repair	0.45	0.15	0.22	461
non functional	0.76	0.67	0.71	2189
accuracy			0.74	5847
macro avg	0.65	0.56	0.58	5847
weighted avg	0.73	0.74	0.72	5847

Overall accuracy score 0.7400376261330597

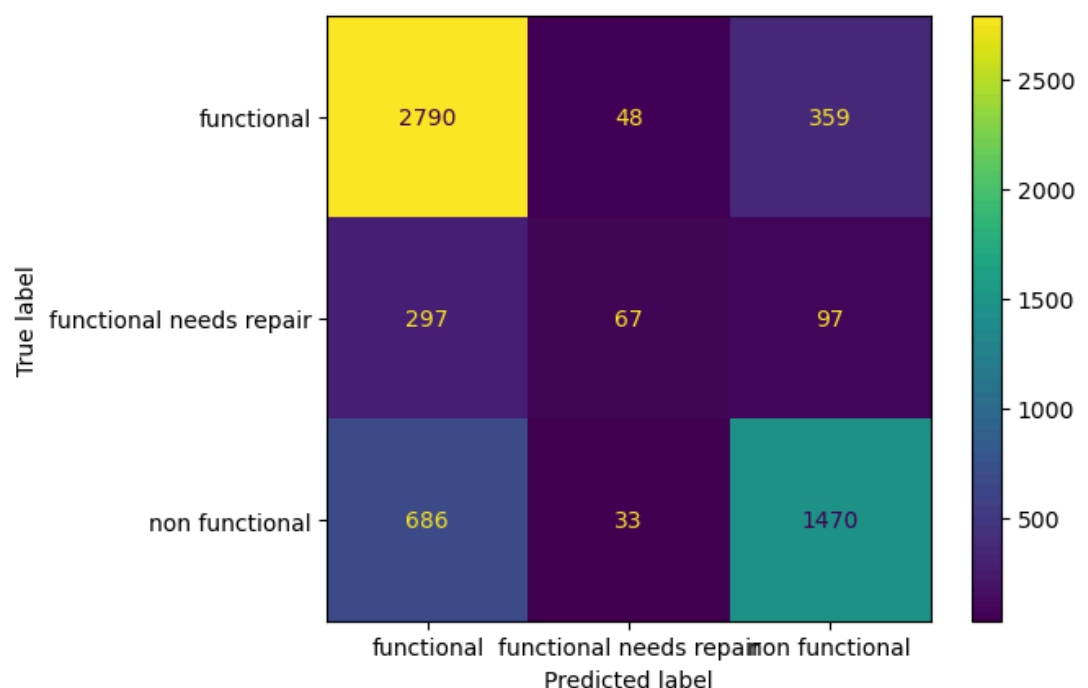
Overall precision score 0.7257561850132295

Overall recall score 0.7400376261330597

Overall F1-score 0.7225622126097963

Confusion Matrix

Out[37]: <sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay at 0x1db80dc3620>







```
In [39]: # we start with a decision tree classifier
# Fit the decision tree classifier to training data
dt_model = pipe_tree.fit(X_train, y_train)

# Print the accuracy on test set
pipe_tree.score(X_test, y_test)
```

Out[39]: 0.7094236360526766

```
In [40]: # Print model metrics
# Predict target
y_pred = dt_model.predict(X_test)

# Create metrics object
score_metrics = Metrics()

# Print classification report, scores, and confusion matrix
score_metrics.model_score(dt_model, y_test, y_pred)
```

	precision	recall	f1-score	support
functional	0.77	0.75	0.76	3197
functional needs repair	0.31	0.30	0.30	461
non functional	0.71	0.73	0.72	2189
accuracy			0.71	5847
macro avg	0.59	0.59	0.59	5847
weighted avg	0.71	0.71	0.71	5847

Overall accuracy score 0.7094236360526766

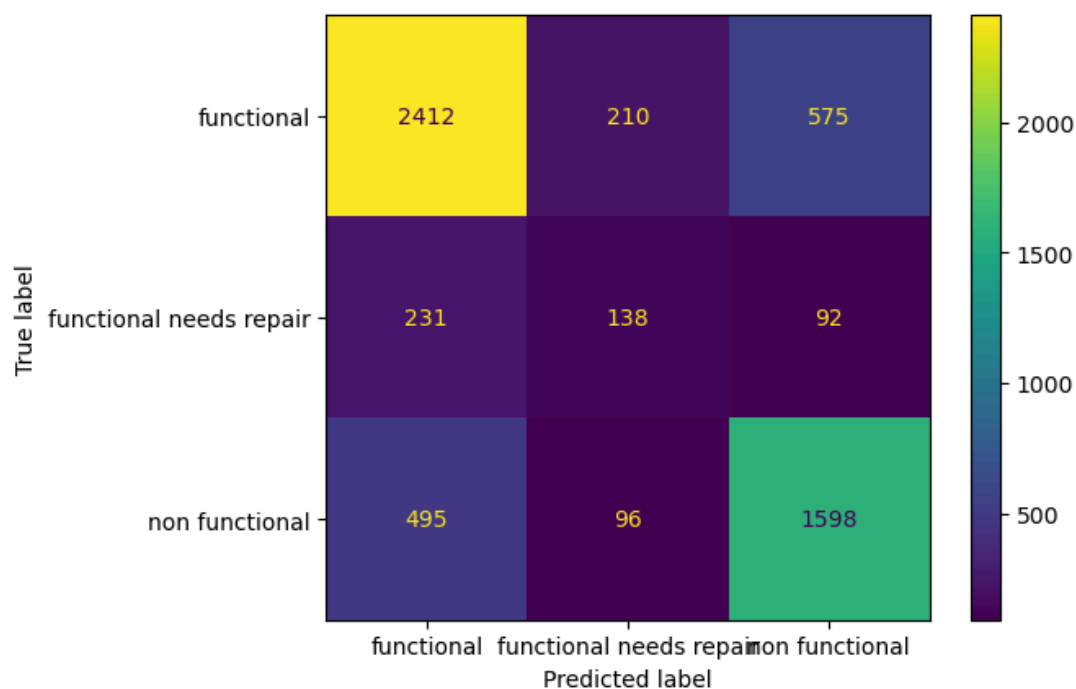
Overall precision score 0.7089129916423678

Overall recall score 0.7094236360526766

Overall F1-score 0.7090454698669827

Confusion Matrix

Out[40]: <sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay at 0x1db819a8ef0>





From the 2 baseline models, we can observe that the Logistic regression has the highest accuracy score of about 0.773 and also the highest f1-score of about 0.766.

The plots of features importances show that the 10 most influential features are:

- i.) Longitude
- ii.) Latitude
- iii.) Quantity\_group
- iv.) gps height
- v.) Population
- vi.) Water type group
- vii.) Construction year
- viii.) Payment type or method
- ix.) Extraction type class
- x.) amount\_tsh

## CROSS VALIDATION MODELS

K-fold cross validation may improve on the performance of training and test splits because it splits the entire dataset into K equal sections of data. At each iteration, it fits the model on different sections of the data and runs a test on the remaining test set, which must also be different for each iteration.

The scores are finally averaged out, which means below-average and above-average score may cancel each other out to give a more realistic score.

```
In [42]: ▶ # Create a Function for Cross Validation test
def cross_val_test(K_folds, model_instance, X, y):

    # Plot a distribution of train and test accuracies over a range of K
    train_accuracy = []
    test_accuracy = []

    for fold in K_folds:
        # Instantiate classifier model
        model = model_instance

        # Perform cross-validation on full dataset to evaluate model performance
        cv_scores = cross_validate(model, X, y, cv=fold, return_train_score=True)

        # Find mean train and test scores
        mean_train_score = np.mean(cv_scores['train_score'])
        mean_test_score = np.mean(cv_scores['test_score'])


        # Append the lists for mean scores
        train_accuracy.append(mean_train_score)
        test_accuracy.append(mean_test_score)

    # Plot the mean train and test scores against K fold
    plt.plot(K_folds, train_accuracy, label='Train score')
    plt.plot(K_folds, test_accuracy, label='Test score')
    plt.xlabel('Cross Validation Fold')
    plt.ylabel('Accuracy Score')
    plt.legend();
    return f"This is the Cross Validation Test for {model_instance}"
```

## LOGISTIC REGRESSION CROSS VALIDAITION

```
In [43]: # Plot a distribution of train and test accuracies over a range of K folds  
#  
cross_val_test(range(5, 18), LogisticRegression(random_state=42), X, y)
```


Out[43]: 'This is the Cross Validation Test for LogisticRegression(random\_state=42)'



## DECISION TREE CLASSIFIER WITH CROSS VALIDATION

```
In [44]: cross_val_test(range(5,18), DecisionTreeClassifier(random_state=42),X ,y
```

```
Out[44]: 'This is the Cross Validation Test for DecisionTreeClassifier(random_s  
tate=42)'
```



From the distribution of accuracy across different K folds, we observe the test accuracy score closely matches that of the train score. Therefore, the model is not overfitting. Both train and test accuracy scores are at their highest and closest at K-fold=6.

However, the accuracy scores are lower than that of the baseline logistic regression model. This shows that cross validation does not improve performance.

## HANDLING CLASS IMBALANCE

```
In [45]: ▶ # Previous original class distribution
print(y_train.value_counts())

# Fit SMOTE to training data
X_train_resampled, y_train_resampled = SMOTE().fit_resample(X_train, y_train)

# Preview synthetic sample class distribution
print('\n')
print(pd.Series(y_train_resampled).value_counts())
```

status_group	
functional	7492
non functional	5202
functional needs repair	949

Name: count, dtype: int64

status_group	
functional	7492
non functional	7492
functional needs repair	7492

Name: count, dtype: int64

```
In [46]: # Print the baseline logistic accuracy and f1 score on test set
y_pred = logistic_model.predict(X_test)
print("Baseline Logistic Regression Accuracy", logistic_model.score(X_test, y_test))
print("Baseline Logistic Regression F1-score", f1_score(y_test, y_pred, average='weighted'))

# Print the baseline decision tree accuracy and f1 score on test set
y_pred = dt_model.predict(X_test)
print("Baseline Decision Tree Accuracy", pipe_tree.score(X_test, y_test))
print("Baseline Decision Tree F1-score", f1_score(y_test, y_pred, average='weighted'))

# Fit the logistic regression model on resampled data
log_model_resampled = pipe_logistic.fit(X_train_resampled, y_train_resampled)
y_pred_log = log_model_resampled.predict(X_test)

# Fit the decision tree classifier on resampled data
dt_model_resampled = pipe_tree.fit(X_train_resampled, y_train_resampled)
y_pred_dt = dt_model_resampled.predict(X_test)

# Print the accuracy and f1 scores on SMOTE sample
print("Logistic Regression Accuracy with SMOTE", log_model_resampled.score(X_test, y_test))
print("Logistic Regression F1-score with SMOTE", f1_score(y_test, y_pred_log, average='weighted'))

print("Decision Tree Accuracy with SMOTE", dt_model_resampled.score(X_test, y_test))
print("Decision Tree F1-score with SMOTE", f1_score(y_test, y_pred_dt, average='weighted'))
```

```
Baseline Logistic Regression Accuracy 0.5500256541816316
Baseline Logistic Regression F1-score 0.46841064174267905
Baseline Decision Tree Accuracy 0.383615529331281
Baseline Decision Tree F1-score 0.4429512056722699
Logistic Regression Accuracy with SMOTE 0.737130152214811
Logistic Regression F1-score with SMOTE 0.7233501724332839
Decision Tree Accuracy with SMOTE 0.7080554130323242
Decision Tree F1-score with SMOTE 0.708735081762939
```

## RANDOM FOREST CLASSIFICATION

### BASELINE RANDOM CLASSIFIER

I will start by instantiating and fitting a Random Forest classifier with default parameters.



```
In [47]:  # Define a pipeline for Random Forest classification
pipe_rf = Pipeline([('ss', StandardScaler()),
                    ('rf', RandomForestClassifier(random_state=42))])

# Fit the Random Forest pipeline on training data
rf_model = pipe_rf.fit(X_train, y_train)

# Training accuracy score
print("Random Forest Accuracy on Training Set", rf_model.score(X_train,

# Test accuracy score
print("Random Forest Accuracy on Test Set", rf_model.score(X_test, y_test))

Random Forest Accuracy on Training Set 0.9998534046763908
Random Forest Accuracy on Test Set 0.7762955361723961
```

```
In [48]:  #Print model metrics
# Predict target
y_pred = rf_model.predict(X_test)

# Create metrics object
score_metrics = Metrics()

# Print model metrics
metrics_list=[]
score_metrics.get_metrics("Random Forest", rf_model, y_test, y_pred)
```

```
Out[48]:  [{'model_name': 'Random Forest',
            'accuracy': 0.7762955361723961,
            'f1 score': 0.7662198899017006,
            'precision': 0.7655272454665789,
            'recall': 0.7762955361723961}]
```

The baseline Random Forest classifier performance is the highest so far at about 0.776 on the test set. However, the training set accuracy is 0.999, which indicates that the model is overfitting.

Next I will perform a cross validation test on the model and note the change in performance.

Cross validation has improved accuracy from about 0.777 to about 0.782 but the model is still overfitting.

## CONCLUSION

From the analytics section, I observed the following:

### Relationship Between Pump Functionality and Continuous Variables

1. For the total static head feature (amount\_tsh), waterpoints with zero static head have the highest density of pumps overall. Also, among the three pump classes at this point, non-functional pumps have the highest density followed by functional pumps. Functional-needs-repair pumps are the least. However, the dataset has a high number

(59.2%) of pumps with zero static head and there is no information about the high occurrence.

Also, while it's expected that a pump with zero static head means that it is faulty, some of these pumps are functional and it's not possible to explain how that comes about.

From the box plot of total\_static\_head vs. pump condition, we can see that the pumps having tsh above approx. 125,000 are all functional. High static head may be an important feature because the higher the tsh the higher the probability of a pump being functional.

2. For the GPS height feature, functional and non-functional pumps are fairly equal when the GPS height is zero, at which point the functional-needs-repair pumps are much fewer.
3. For the population feature, waterpoints located in areas with zero population have the highest density of non-functional pumps. This is followed by functional pumps and functional-needs-repair pumps are the least. However, there is no information about whether the wells have been abandoned or the population has relocated.
4. A KDE (kernel density estimation) plot shows that the density of functional pumps is higher among the newest pumps while non-functional pumps are higher among the

From the predictive section, I conclude that it's possible to correctly **predict at atleast 80% accuracy**, the condition of a pump given the data features from the Ministry of Water in Tanzania.

It has an **accuracy score of 0.800 (80%)**, an **F1-score of 0.788**, a **precision of 0.797**, and a **recall of 0.800**. Even though it has slightly higher scores, the train scores show that it is overfitting the data while the XGBoost is a good fit (not underfitting or overfitting).

To save on the cost of attending to pumps whose condition has been erroneously predicted, or ignoring faulty pumps predicted to be in good condition, I started with the intention of minimizing both False Positives (high precision) and False Negatives (high recall). The XGBoost model has high overall scores for both precision and recall, at 0.797 and 0.8 respectively.

The classification report also shows that the XGBoost model has high scores for the functional class (precision=0.78, recall=0.91, f1-score=0.84) as well as the non-functional class (precision=0.85, recall=0.74, f1-score=0.79) but the score for the functional-needs-repair class are low (precision=0.62, recall=0.21, f1-score=0.31). This could be attributed to the fact that there were only 4.42% (2,642 out 59,400) counts of these class in the dataset. This means the model did not learn about this model as much as it did the other two classes.

Some of the top features influencing a prediction include:

- i.) quantity-group (the quantity of water)
- ii.) The water point type
- iii.) The extraction type class
- iv.) The basin
- v.) scheme management
- vi.) The installer
- vii.) payment type

